

COMP348 — Document Processing and the Semantic Web

Week 02 Lecture 1: Searching for Information

Diego Mollá

Department of Computer Science
Macquarie University

COMP348 2019H1

Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Reading

Essential Reading

- NLTK chapter 6 section 3.3 (precision and recall).
- Manning et al. IR book, chapter 1 (Boolean retrieval), chapter 6 section 2 (Td.idf), chapter 8 section 3 (precision and recall). <http://nlp.stanford.edu/IR-book/>

Additional Reading

- Brin and Page (1998):
<http://infolab.stanford.edu/~backrub/google.html> — a famous paper by the founders of Google.

Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Need for Search

The Problem

- The Web can be seen as a very large, unstructured data store.
- There exist hundreds of millions of Web pages but there is no central index.
- Even worse: It is unknown where all the Web servers are.

The Solution

Search engines.

Information Retrieval

Information Retrieval (IR)

- IR is about searching for information.
- IR typically means “document retrieval”.
- IR is one of the core components of Web search.



<http://boston.lti.cs.cmu.edu/classes/11-744/treclogo-c.gif>

Stages in an IR System

1: Indexing

- This stage is done off-line, prior to any searches.
- The goal is to reduce the documents to a description: the indices.
- Optimise the representation: ignore the terms that do not contribute.

2: Retrieval

- Use the indices to retrieve the documents (ignore the remaining information in the documents).

Programme

- 1 Information Retrieval
- 2 **Evaluation**
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Why Evaluate?

- Document processing systems almost never give 100% correct results.
- When you develop a document processing system, you want to know how good it is.
- You want to know if a modification in a system is an improvement.
- Human evaluations are expensive to produce.
- In this lecture we will focus on automatic evaluations.

Of course, **in addition** you have to debug the system.

Training vs. Test Data

- For pretty much all evaluations, you want to divide your data into at least two sets: training and test.
- **Training** data is what you use to develop your models.
 - You only look at the training data.
 - For statistical models (coming later in this course), this is what you use to calculate your statistics.
- **Test** data is separate.
- You may also have a third set of data to help develop your system (DevTest).
 - You'll see the use of the DevTest set when we look at statistical models.

Golden Rule

You don't ever, ever, look at the test data
(you only look at its evaluation results).

Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Types of Errors

Errors by a system making a binary choice can often be broken into two types:

- 1 Selecting something when it's not supposed to be selected.
- 2 Not selecting something when it is supposed to be selected.

Examples

- 1 If the task is to identify whether an email is spam or not, the system can mistakenly classify an email as spam when it is not spam.
- 2 If the task is to identify documents as relevant or not, the system can mistakenly classify irrelevant documents as relevant, or relevant documents as irrelevant.

Positives and Negatives

We can group results of the system into four categories: **tp** (true positive), **fp** (false positive), **fn** (false negative), **tn** (true negative)

| system decision | actual case | |
|--------------------|-------------|------------|
| | target | not target |
| selected | tp | fp |
| not selected | fn | tn |

Example: Positives and Negatives in Information Retrieval

In IR, “relevant” documents belong to the target category.

| system decision | actual case | |
|--------------------|-------------|--------------|
| | relevant | not relevant |
| retrieved | tp | fp |
| not retrieved | fn | tn |

- Our retrieval system fails to retrieve a relevant document: this is a **false negative**.

Example: Spam Filtering

In spam filtering, “spam” emails belong to the target category.

| system | actual case | |
|-----------------|-------------|----------|
| | spam | not spam |
| marked spam | tp | fp |
| not marked spam | fn | tn |

- Our spam filter classifies a legitimate email as spam: this is a **false positive**.

Question

False positives in spam filtering are usually more dangerous than false negatives; why?

Example: Spam Filtering

In spam filtering, “spam” emails belong to the target category.

| system | actual case | |
|-----------------|-------------|----------|
| | spam | not spam |
| marked spam | tp | fp |
| not marked spam | fn | tn |

- Our spam filter classifies a legitimate email as spam: this is a **false positive**.

Question

False positives in spam filtering are usually more dangerous than false negatives; why?

Precision and Recall

Formulas

- $\text{precision} = \text{tp} / (\text{tp} + \text{fp})$
- $\text{recall} = \text{tp} / (\text{tp} + \text{fn})$

Example

From a total collection of 200 documents, a retrieval system returned 30 documents, but 5 were not relevant. It also missed 12 documents.

Example

| system | actual case | |
|--------------|-------------|------------|
| | target | non target |
| selected | 25 | 5 |
| not selected | 12 | 158 |

Values of measures

- precision = $25/30$
- recall = $25/37$

Accuracy

- Accuracy is the number correctly classified out of the whole set.
 - $\text{accuracy} = (tp + tn) / (tp + fp + tn + fn)$
 - For previous example, accuracy is 183/200
- Sometimes used (inaccurately) to refer to precision.

Question

What happens if you have unbalanced classes (e.g. 90% of the data belongs to class 1)?

F-Measure

- Another way of getting a single measure for a system is to combine precision and recall.
- For the general case,

$$F_{\beta} = (1 + \beta^2) \frac{\text{precision} \times \text{recall}}{\beta^2 \text{precision} + \text{recall}}$$

- The most commonly used instance is when $\beta = 1$, referred to as F_1 :

$$F_1 = 2 \frac{\text{precision} \times \text{recall}}{\text{precision} + \text{recall}}$$

- This is just the **harmonic mean** of precision and recall.
- For previous example, $F_1 = 0.746$

Exercise: Spam Filtering

Exercise

Assume your system processes 1000 emails. It classifies 640 as spam, of which 480 are actually spam. It missed 120 spam emails. What are the precision and recall of the spam detection and the accuracy of the system?

Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval**
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Bag of Words Representation

Bag of words (BoW)

- At indexing time, a compact representation of the document is built.
- The document is seen as a bag of words.
- Information about word position is (often) discarded.
- Only the important words are kept.

The bag-of-words model is a simplifying representation used in natural language processing and information retrieval (IR). In this model, a text (such as a sentence or a document) is represented as the bag (multiset) of its words, disregarding grammar and even word order but keeping multiplicity. Recently, the bag-of-words model has also been used for computer vision.



{bag, bag-of-words, computer, disregarding, document, grammar, information, IR, keeping, language, model, multiplicity, multiset, natural, order, processing, representation, represented, retrieval, sentence, simplifying, text, vision, word, words}

Stop Words

Stop words

- A simple solution to determine important words is to keep a list of non-important words: the **stop words**.
- All stop words in a document are ignored.
- Stop words are language-specific.
- Typically, stop words are connecting words.

Stop words in NLTK

```
>>> from nltk.corpus import stopwords
>>> stop = stopwords.words('english')
>>> stop[:5]
['i', 'me', 'my', 'myself', 'we']
```

Term Frequency

Term Frequency

- Words that are not frequent are usually not important.
- Words that are too frequent may occur in most documents and therefore can't be used to discriminate among documents.
- Usually, important words are in the middle.

Zipf's Law for term frequency


- A small percentage of words are very frequent.
- A large percentage of words have very little frequency.
- The relation approximates a Zipfian distribution.
- This is also referred as “long-tailed” distribution.

Zipf's Law in Action

Python code

```
import nltk
import collections
import matplotlib.pyplot as plt
words = nltk.corpus.gutenberg.words('austen-emma.txt')
fd = collections.Counter(words)
data = sorted([fd[k] for k in fd], reverse=True)
plt.plot(data[:1000])
plt.show()
```

500 most frequent words



tf.idf

tf.idf

- **Term frequency:** If a word is very frequent in a document, it is important for the document.

$$tf(t, d) = \text{frequency of word } t \text{ in document } d$$

- **Inverse document frequency:** If a word appears in many documents, it is not important for any document.

$$idf(t) = \log \frac{\text{number of documents}}{\text{number of documents that contain } t}$$

- *tf.idf* combines these two characteristics.

$$tf.idf(t, d) = tf(t, d) \times idf(t)$$

Problems with Bag of Word Representations

BoW representations ignore important information such as:

Word position: “Australia beat New Zealand” is not the same as “New Zealand beat Australia”

Morphology: If you search for “table”, a webpage that uses the word “tables” might be relevant.

Words with similar meanings: If you search for “truck”, a webpage that uses the word “lorry” might be relevant.

Ambiguity: If you search for “Apple” you might be interested in the company and not in the fruit.

Still, BoW representations are very simple, fast, and often surprisingly good.

Beyond BoW Representations

- A simple way to account for (some) information about word positions is to use **n-grams**:
 - Bigrams, trigrams, 4-grams (usually there is no need for longer n-grams).
- Thus, instead of representing a text as a bag of words, it can be represented as a bag of n-grams.
- Using n-grams instead of words may introduce other kinds of problems (we will see some of these problems in a future lecture).

Accounting for Word Meaning I


Ambiguity

- Word disambiguation attempts to determine the sense of a word.
- A word like “Apple” could be disambiguated as “apple1” or “apple2” to account for its several meanings.
- Word disambiguation systems usually look at the “context” of the word:
 - Yesterday I ate an apple₁.
 - Apple₂ reported a benefit last fiscal year.

Accounting for Word Meaning II

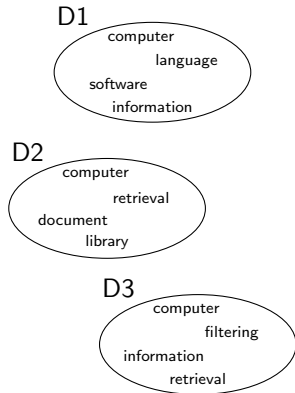
Synonymy

- There are lexical resources such as thesauri (singular: thesaurus) that list words with related meanings.
 - WordNet is a popular resource (<https://wordnet.princeton.edu/>)
- Recent innovations include the use of **distributional semantics** to determine the similarity between two words.
 - Word2Vec and Glove are two systems that map words to vectors called **word embeddings**.

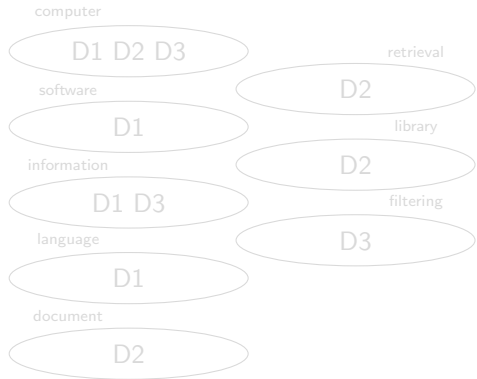


Inverted Indices

Index




Inverted Index




Inverted Indices

Index



Inverted Index



Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - **Boolean Retrieval**
 - Vector Retrieval
 - Vector Retrieval in Python

Retrieval

- In the retrieval stage, the index is searched.
- This enables fast retrieval.
- Note that the index does not contain the full information from the documents.
- For example, searching a stop word will be useless.

Boolean Retrieval

- Use Boolean operations among the search terms.
 - x AND y Documents that contain both terms.
 - x OR y Documents that contain at least one term.
 - NOT x Documents that do not contain the term.
- The use of inverted indices simplifies this method.
 - x AND y Set intersection.
 - x OR y Set union.
 - NOT x Set complement.

Example of Boolean Retrieval

Keywords

D1: {computer, software, information, language}

D2: {computer, document, retrieval, library}

D3: {computer, information, filtering, retrieval}

Inverted Index

computer \rightarrow {D1, D2, D3}, software \rightarrow {D1}, information \rightarrow {D1, D3},

language \rightarrow {D1}, document \rightarrow {D2}, retrieval \rightarrow {D2, D3},

library \rightarrow {D2}, filtering \rightarrow {D3}

Boolean Query

(information OR document) AND retrieval

Result

$$(\{D1, D3\} \cup \{D2\}) \cap \{D2, D3\} = \{D2, D3\}$$


Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - **Vector Retrieval**
 - Vector Retrieval in Python

Vector Retrieval

Boolean retrieval and ranking

- There are no obvious methods to rank the results of Boolean retrieval.
 - Google introduced PageRank but we will see this later...
- An easy method to rank documents is to represent them as vectors and use well-established methods for vector comparison.



Vector Space Model

Template:

{computer,software,information,document,retrieval,language,library,filtering}

Initial documents

D1:{computer, software, information, language}

D2:{computer, document, retrieval, library}

D3:{computer, information, filtering, retrieval}

Document vectors

D1: (1,1,1,0,0,1,0,0)

D2: (1,0,0,1,1,0,1,0)

D3: (1,0,1,0,1,0,0,1)

Document matrix

(typically a **sparse matrix**)

$$D = \begin{pmatrix} 1 & 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 1 & 1 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 1 & 0 & 0 & 1 \end{pmatrix}$$


Information in the Vectors

- In the vector space model, each word in the vocabulary represents an element in a vector.
- The final document matrix will typically be **sparse** since a document will typically contain only a small fraction of all the possible words.
- Possible information to store in the vector:
 - The occurrence of a word/stem/n-gram (1) or not (0) \leftarrow as in our example.
 - The word frequency.
 - *tf.idf* \leftarrow a popular choice.
 - Distributional semantics \leftarrow a hot research topic.

Cosine Similarity

Cosine Method

- Compare the cosine of the angle between vectors.
- If the angle is zero, then the cosine is 1.



$$\begin{aligned}\cos(D_1, Q_1) &= \cos(\alpha) \\ \cos(D_2, Q_1) &= \cos(0) = 1 \\ \cos(D_3, Q_1) &= \cos(\beta)\end{aligned}$$

Cosine Similarity: Formulas

General Formula

$$\cos(D_j, Q_k) = \frac{\sum_{i=1}^N D_{j,i} Q_{k,i}}{\sqrt{\sum_{i=1}^N D_{j,i}^2} \sqrt{\sum_{i=1}^N Q_{k,i}^2}} = \frac{D_j \cdot Q_k}{\|D_j\|_2 \|Q_k\|_2}$$

If the vectors are normalised

$$\cos(D_j, Q_k) = \sum_{i=1}^N D_{j,i} Q_{k,i} = D_j \cdot Q_k$$


Programme

- 1 Information Retrieval
- 2 Evaluation
 - Precision and Recall
- 3 Indexing and Retrieval
 - Indexing
 - Boolean Retrieval
 - Vector Retrieval
 - Vector Retrieval in Python

Vectors and Matrices in Python

numpy

- Python's numpy is a collection of libraries that include manipulation of vectors and matrices.
- <http://www.numpy.org/>
- It's pre-loaded in the Anaconda distribution.



Manipulating Vectors

```
>>> import numpy as np
>>> a = np.array([1,2,3,4])
>>> a[0]
1
>>> a[1:3]      # slicing
array([2, 3])
>>> a+1         # add a constant to a vector
array([2, 3, 4, 5])
>>> b=np.array([2,3,4,5])
>>> a+b         # add two vectors
array([3, 5, 7, 9])
>>> a*b         # pairwise multiplication
array([ 2,  6, 12, 20])
>>> np.dot(a,b) # dot product between vectors, a . b
40
```

Manipulating Matrices

```
>>> x = np.array([[1,2,3],[4,5,6]])
>>> x
array([[1, 2, 3],
       [4, 5, 6]])
>>> y = np.array([[1,1,1],[2,2,2]])
>>> x+y           # add two matrices
array([[2, 3, 4],
       [6, 7, 8]])
>>> x*y           # pairwise multiplication
array([[ 1,  2,  3],
       [ 8, 10, 12]])
>>> x.T           # transpose
array([[1, 4],
       [2, 5],
       [3, 6]])
>>> np.dot(x.T,y) # dot product
array([[ 9,  9,  9],
       [12, 12, 12]])
```


Scikit-learn I

- <http://scikit-learn.org/>
- Incorporates an extensive set of machine learning algorithms into Python.
- It has a consistent and intuitive interface.
- The documentation is very complete.
- Includes generic tutorials on the main machine learning algorithms.
- It's pre-loaded in the Anaconda distribution.

Scikit-learn II

The screenshot shows the scikit-learn.org website. At the top, there's a navigation bar with links: Home, Installation, Documentation, and Examples. A Google Custom Search bar is also present. The main header features the scikit-learn logo and the tagline "Machine Learning in Python". Below this, a grid of 12 small images represents different machine learning topics. To the right of the grid, a list of bullet points highlights key features: simple and efficient tools, accessibility, built on NumPy, SciPy, and matplotlib, and open source/BSD license. Below the main header, there are four columns, each representing a machine learning task: Classification, Regression, Clustering, Dimensionality reduction, Model selection, and Preprocessing. Each column contains a brief description, applications, and algorithms.

scikit-learn
Machine Learning in Python

- Simple and efficient tools for data mining and data analysis
- Accessible to everybody, and reusable in various contexts
- Built on NumPy, SciPy, and matplotlib
- Open source, commercially usable - BSD license

Classification

Identifying to which set of categories a new observation belong to.

Applications: Spam detection, Image recognition.

Algorithms: SVM, nearest neighbors, random forest, ...

Regression

Predicting a continuous value for a new example.

Applications: Drug response, Stock prices.

Algorithms: SVR, ridge regression, Lasso, ...

Clustering

Automatic grouping of similar objects into sets.

Applications: Customer segmentation, Grouping experiment outcomes

Algorithms: k-Means, spectral clustering, mean-shift, ...

Dimensionality reduction

Reducing the number of random variables to consider.

Applications: Visualization, Increased efficiency

Algorithms: PCA, Isomap, non-negative matrix factorization.

Model selection

Comparing, validating and choosing parameters and models.

Goal: Improved accuracy via parameter tuning

Modules: grid search, cross validation, metrics.

Preprocessing

Feature extraction and normalization.

Application: Transforming input data such as text for use with machine learning algorithms.

Modules: preprocessing, feature extraction.

tf.idf with scikit-learn

```
>>> import glob
>>> files = glob.glob('enron1/ham/*.txt')
>>> from sklearn.feature_extraction.text import TfidfVectorizer
>>> tfidf = TfidfVectorizer(input='filename', stop_words='eng')
>>> tfidf_values = tfidf.fit_transform(files)
>>> len(tfidf.get_feature_names())
19892
>>> tfidf.get_feature_names()[10000:10005]
['grandma', 'grandpa', 'grandsn', 'grandsons', 'grant']
>>> type(tfidf_values)
scipy.sparse.csr.csr_matrix
>>> type(tfidf_values.toarray())
numpy.ndarray
>>> tfidf_values.shape
(3672, 19892)
```

Normalised tf.idf and cosine similarity in Python

```
>>> tfidf_norm = TfidfVectorizer(input='filename',  
                                stop_words='english',  
                                norm='l2')  
>>> tfidf_norm_values = tfidf_norm.fit_transform(files).toarray()  
>>> import numpy as np  
>>> def cosine_similarity(X,Y):  
    return np.dot(X,Y)  
>>> cosine_similarity(tfidf_norm_values[0,:],  
                    tfidf_norm_values[1,:])  
0.017317648885111028
```

Some Open Source Search Engines

If you don't want to implement your search engine from scratch, try these (<http://www.mytechlogy.com/IT-blogs/8685/top-5-open-source-search-engines/>):

- 1 Elasticsearch: <https://www.elastic.co/>
- 2 Solr: <https://lucene.apache.org/solr/>
- 3 Lucene: <https://lucene.apache.org/>
- 4 Sphinx: <http://sphinxsearch.com/>
- 5 Xapian: <https://xapian.org/>
- 6 Indri: <https://www.lemurproject.org/>
- 7 Zettair: <http://www.seg.rmit.edu.au/zettair/>

The following is a Python library that can be used for indexing and retrieving documents (among many other things):

- 1 Gensim: <https://radimrehurek.com/gensim/>

Take-home Messages

- 1 What is indexing? what is retrieval?
- 2 What is an inverted index?
- 3 Perform Boolean retrieval by hand.
- 4 Implement Boolean retrieval in Python.
- 5 Use sklearn to build a vector model with tf.idf.
- 6 Use sklearn to implement cosine similarity.

What's Next

Friday

- Web Search

Reading

- Brin and Page (1998) — a famous paper by the founders of Google.